



# Some underlying foundations for RL

Dale Schuurmans


Google Brain



# The RL problem



# The RL problem



1. multi-agent interaction → non-stationarity
2. partial observability → construct memory
3. exploration → explore/exploit
4. sequential decisions → temporal credit assignment
5. exploitation → policy optimization

# Optimizing one step decision making

Batch policy optimization

# Batch policy optimization

Given data

	$a_1$	$a_2$	$\dots$	$a_n$
$x_1$		$r_1$		
$x_2$				$r_2$
$x_3$	$r_3$			
$x_4$			$r_4$	
$x_5$	$r_5$			
$x_6$		$r_6$		
$\vdots$			$r_i$	
$x_m$				$r_m$

Optimize policy  $\pi : X \rightarrow \Delta^n$  to maximize expected reward on **test** contexts

$$\pi(a \mid x) = e^{q(x)_a - F(q(x))}$$

$$F(q(x)) = \log \sum_a e^{q(x)_a}$$

$$q : X \rightarrow \mathbb{R}^n \quad \text{neural network}$$

# Batch policy optimization

Given data

	$a_1$	$a_2$	...	$a_n$
$x_1$		$r_1$		
$x_2$				$r_2$
$x_3$	$r_3$			
$x_4$			$r_4$	
$x_5$	$r_5$			
$x_6$		$r_6$		
$\vdots$			$r_i$	
$x_m$				$r_m$

## Three key issues

1. generalization
2. optimization
3. missing data

Optimize policy  $\pi : X \rightarrow \Delta^n$

$$\pi(a | x) = e^{q(x)_a - F(q(x))}$$

$$F(q(x)) = \log \sum_a e^{q(x)_a}$$

$$q : X \rightarrow \mathbb{R}^n \quad \text{neural network}$$

to maximize expected reward on **test** contexts

# Batch policy optimization

Given data

	$a_1$	$a_2$	...	$a_n$
$x_1$		$r_1, \beta_1$		
$x_2$				$r_2, \beta_2$
$x_3$	$r_3, \beta_3$			
$x_4$			$r_4, \beta_4$	
$x_5$	$r_5, \beta_5$			
$x_6$		$r_6, \beta_6$		
$\vdots$			$r_{\cdot}, \beta_{\cdot}$	
$x_m$				$r_m, \beta_m$

Isn't this a solved problem?

We know how to do this, right?

**Default answer**

- maximize importance corrected expected reward
- (assume have proposal probabilities)

$$\max \sum_i \frac{\pi(a_i | x_i)}{\beta_i} r_i$$

Optimize policy  $\pi : X \rightarrow \Delta^n$  to maximize expected reward on **test** contexts

$$\pi(a | x) = e^{q(x)_a - F(q(x))}$$

$$F(q(x)) = \log \sum_a e^{q(x)_a}$$

$$q : X \rightarrow \mathbb{R}^n \quad \text{neural network}$$

# Batch policy optimization

Given data

	$a_1$	$a_2$	...	$a_n$
$x_1$		$r_1$		
$x_2$				$r_2$
$x_3$	$r_3$			
$x_4$			$r_4$	
$x_5$	$r_5$			
$x_6$		$r_6$		
$\vdots$			$r_i$	
$x_m$				$r_m$

## Three key issues

1. generalization

2. optimization

3. missing data

Importance corrected  
expected reward

okay



Optimize policy  $\pi : X \rightarrow \Delta^n$   
 $\pi(a | x) = e^{q(x)_a - F(q(x))}$

$$F(q(x)) = \log \sum_a e^{q(x)_a}$$

$q : X \rightarrow \mathbb{R}^n$  neural network

to maximize expected reward on **test** contexts



# Optimization objectives

Given data

	$a_1$	$a_2$	$\dots$	$a_n$
$x_1$		$r_1$		
$x_2$				$r_2$
$x_3$	$r_3$			
$x_4$			$r_4$	
$x_5$	$r_5$			
$x_6$		$r_6$		
$\vdots$			$r_i$	
$x_m$				$r_m$

Optimize policy  $\pi : X \rightarrow \Delta^n$  to maximize expected reward on **test** contexts

$$\pi(a | x) = e^{q(x)_a - F(q(x))}$$

$$F(q(x)) = \log \sum_a e^{q(x)_a}$$

$$q : X \rightarrow \mathbb{R}^n \quad \text{neural network}$$

# Optimization objectives

Now assume given **complete** data

	$a_1$	$a_2$	$\dots$	$a_n$
$x_1$	$r_{11}$	$r_{12}$	$r_{1\dots}$	$r_{1n}$
$x_2$	$r_{21}$	$r_{22}$	$r_{2\dots}$	$r_{2n}$
$x_3$	$r_{31}$	$r_{32}$	$r_{3\dots}$	$r_{3n}$
$x_4$	$r_{41}$	$r_{42}$	$r_{4\dots}$	$r_{4n}$
$x_5$	$r_{51}$	$r_{52}$	$r_{5\dots}$	$r_{5n}$
$x_6$	$r_{61}$	$r_{62}$	$r_{6\dots}$	$r_{6n}$
$\vdots$	$r_{:1}$	$r_{:2}$	$r_{: \dots}$	$r_{:n}$
$x_m$	$r_{m1}$	$r_{m2}$	$r_{m\dots}$	$r_{mn}$

Optimize policy  $\pi : X \rightarrow \Delta^n$

$$\pi(a | x) = e^{q(x)_a - F(q(x))}$$

$$F(q(x)) = \log \sum_a e^{q(x)_a}$$

$q : X \rightarrow \mathfrak{R}^n$  neural network

to maximize expected reward on **test** contexts

# Optimization objectives

Now assume given **complete** data

	$a_1$	$a_2$	$\dots$	$a_n$
$x_1$	$r_{11}$	$r_{12}$	$r_{1\dots}$	$r_{1n}$
$x_2$	$r_{21}$	$r_{22}$	$r_{2\dots}$	$r_{2n}$
$x_3$	$r_{31}$	$r_{32}$	$r_{3\dots}$	$r_{3n}$
$x_4$	$r_{41}$	$r_{42}$	$r_{4\dots}$	$r_{4n}$
$x_5$	$r_{51}$	$r_{52}$	$r_{5\dots}$	$r_{5n}$
$x_6$	$r_{61}$	$r_{62}$	$r_{6\dots}$	$r_{6n}$
$\vdots$	$r_{:1}$	$r_{:2}$	$r_{: \dots}$	$r_{:n}$
$x_m$	$r_{m1}$	$r_{m2}$	$r_{m\dots}$	$r_{mn}$

Optimize policy  $\pi : X \rightarrow \Delta^n$   
 $\pi(a | x) = e^{q(x)_a - F(q(x))}$   
 $F(q(x)) = \log \sum_a e^{q(x)_a}$   
 $q : X \rightarrow \mathbb{R}^n$  neural network

## Target objective

- expected reward:  $\max \sum_i \mathbf{r}_i \cdot \pi(x_i)$

Done, right?

Not so fast ...

## This objective has serious problems

- actually trying to solve:  $\max \sum_i \mathbf{r}_i \cdot \mathbf{f}(q(x_i))$
- plateaus everywhere
- can have **exponentially many** local maxima
- nearly impossible to reach a global optima

**Also: you already know not to train this way!**

to maximize expected reward on **test** contexts

# Optimization objectives

Special case: **supervised classification**

	$a_1$	$a_2$	$\dots$	$a_n$
$x_1$	0	1	0	0
$x_2$	0	0	0	1
$x_3$	1	0	0	0
$x_4$	0	0	1	0
$x_5$	1	0	0	0
$x_6$	0	1	0	0
$\vdots$	0	0	1	0
$x_m$	0	0	0	1

Optimize policy  $\pi : X \rightarrow \Delta^n$  to maximize expected **accuracy** on **test** contexts

$$\pi(a | x) = e^{q(x)_a - F(q(x))}$$

$$F(q(x)) = \log \sum_a e^{q(x)_a}$$

$$q : X \rightarrow \mathfrak{R}^n \quad \text{neural network}$$

# Optimization objectives

Special case: **supervised classification**

	$a_1$	$a_2$	$\dots$	$a_n$
$x_1$	0	1	0	0
$x_2$	0	0	0	1
$x_3$	1	0	0	0
$x_4$	0	0	1	0
$x_5$	1	0	0	0
$x_6$	0	1	0	0
$\vdots$	0	0	1	0
$x_m$	0	0	0	1

Optimize policy  $\pi : X \rightarrow \Delta^n$

$$\pi(a | x) = e^{q(x)_a - F(q(x))}$$

$$F(q(x)) = \log \sum_a e^{q(x)_a}$$

$q : X \rightarrow \mathbb{R}^n$  neural network

## Target objective

- expected **accuracy**:  $\max \sum_i \mathbf{r}_i \cdot \pi(x_i)$

But you have never trained with this objective  
Instead, you used a **surrogate objective**

**maximum likelihood**

$$\max \sum_i \mathbf{r}_i \cdot \log \pi(x_i)$$

## What's going on?

- $\mathbf{r}_i \cdot \pi(x_i)$  is differentiable, that's not the issue
- training with  $\mathbf{r}_i \cdot \log \pi(x_i)$  actually achieves better values of  $\mathbf{r}_i \cdot \pi(x_i)$  on the training data

to maximize expected **accuracy** on **test** contexts

# Optimization objectives

Special case: **supervised classification**

	$a_1$	$a_2$	$\dots$	$a_n$
$x_1$	0	1	0	0
$x_2$	0	0	0	1
$x_3$	1	0	0	0
$x_4$	0	0	1	0
$x_5$	1	0	0	0
$x_6$	0	1	0	0
$\vdots$	0	0	1	0
$x_m$	0	0	0	1

**Why?**

- expected accuracy:  $\max \sum_i \mathbf{r}_i \cdot \pi(x_i)$
- maximum likelihood:  $\max \sum_i \mathbf{r}_i \cdot \log \pi(x_i)$

**Useful properties of maximum likelihood**

- $\mathbf{r}_i \cdot \log \pi(x_i)$  is **concave** in  $\mathbf{q}(x_i)$
- it is also **calibrated** w.r.t.  $\mathbf{r}_i \cdot \pi(x_i)$ :

$$\forall \epsilon > 0 \exists \delta > 0 \quad \mathbf{r} \cdot \log \pi^* - \mathbf{r} \cdot \log \pi < \delta \Rightarrow \mathbf{r} \cdot \pi^* - \mathbf{r} \cdot \pi < \epsilon$$

Optimize policy  $\pi : X \rightarrow \Delta^n$

$$\pi(a | x) = e^{q(x)_a - F(q(x))}$$



$$F(q(x)) = \log \sum_a e^{q(x)_a}$$

$q : X \rightarrow \mathbb{R}^n$  neural network

to maximize expected **accuracy** on **test** contexts

# Optimization objectives

Misclassification error on MNIST training data



# Optimization objectives

Back to **general** rewards

	$a_1$	$a_2$	$\dots$	$a_n$
$x_1$	$r_{11}$	$r_{12}$	$r_{1\dots}$	$r_{1n}$
$x_2$	$r_{21}$	$r_{22}$	$r_{2\dots}$	$r_{2n}$
$x_3$	$r_{31}$	$r_{32}$	$r_{3\dots}$	$r_{3n}$
$x_4$	$r_{41}$	$r_{42}$	$r_{4\dots}$	$r_{4n}$
$x_5$	$r_{51}$	$r_{52}$	$r_{5\dots}$	$r_{5n}$
$x_6$	$r_{61}$	$r_{62}$	$r_{6\dots}$	$r_{6n}$
$\vdots$	$r_{:1}$	$r_{:2}$	$r_{:\dots}$	$r_{:n}$
$x_m$	$r_{m1}$	$r_{m2}$	$r_{m\dots}$	$r_{mn}$

Optimize policy  $\pi : X \rightarrow \Delta^n$  to maximize expected reward on **test** contexts

$$\pi(a | x) = e^{q(x)_a - F(q(x))}$$

$$F(q(x)) = \log \sum_a e^{q(x)_a}$$

$$q : X \rightarrow \mathfrak{R}^n \quad \text{neural network}$$



# Optimization objectives

Back to **general** rewards

	$a_1$	$a_2$	...	$a_n$
$x_1$	$r_{11}$	$r_{12}$	$r_{1...}$	$r_{1n}$
$x_2$	$r_{21}$	$r_{22}$	$r_{2...}$	$r_{2n}$
$x_3$	$r_{31}$	$r_{32}$	$r_{3...}$	$r_{3n}$
$x_4$	$r_{41}$	$r_{42}$	$r_{4...}$	$r_{4n}$
$x_5$	$r_{51}$	$r_{52}$	$r_{5...}$	$r_{5n}$
$x_6$	$r_{61}$	$r_{62}$	$r_{6...}$	$r_{6n}$
$\vdots$	$r_{:1}$	$r_{:2}$	$r_{:...}$	$r_{:n}$
$x_m$	$r_{m1}$	$r_{m2}$	$r_{m...}$	$r_{mn}$

## Target objective

- expected reward:  $\max \sum_i \mathbf{r}_i \cdot \boldsymbol{\pi}(x_i)$   
“cost sensitive classification”

Calibrated surrogates exist for  $\mathbf{r}_i \cdot \boldsymbol{\pi}(x_i)$   
(Pires et al. ICML-2013)

## Interesting alternative

- entropy regularized expected reward  
 $\max \sum_i \mathbf{r}_i \cdot \boldsymbol{\pi}(x_i) - \tau \boldsymbol{\pi}(x_i) \cdot \log \boldsymbol{\pi}(x_i)$

Optimize policy  $\pi : X \rightarrow \Delta^n$   
 $\pi(a | x) = e^{q(x)_a - F(q(x))}$

$$F(q(x)) = \log \sum_a e^{q(x)_a}$$

$q : X \rightarrow \mathbb{R}^n$  neural network

to maximize expected reward on **test** contexts

# Optimization objectives

Back to **general** rewards

	$a_1$	$a_2$	$\dots$	$a_n$
$x_1$	$r_{11}$	$r_{12}$	$r_{1\dots}$	$r_{1n}$
$x_2$	$r_{21}$	$r_{22}$	$r_{2\dots}$	$r_{2n}$
$x_3$	$r_{31}$	$r_{32}$	$r_{3\dots}$	$r_{3n}$
$x_4$	$r_{41}$	$r_{42}$	$r_{4\dots}$	$r_{4n}$
$x_5$	$r_{51}$	$r_{52}$	$r_{5\dots}$	$r_{5n}$
$x_6$	$r_{61}$	$r_{62}$	$r_{6\dots}$	$r_{6n}$
$\vdots$	$r_{:1}$	$r_{:2}$	$r_{: \dots}$	$r_{:n}$
$x_m$	$r_{m1}$	$r_{m2}$	$r_{m\dots}$	$r_{mn}$

**Entropy regularized expected reward**

$$\begin{aligned} & \arg \max \mathbf{r} \cdot \boldsymbol{\pi} - \tau \boldsymbol{\pi} \cdot \log \boldsymbol{\pi} \\ &= \arg \min \tau F(\mathbf{r}/\tau) - \mathbf{r} \cdot \boldsymbol{\pi} + \tau F^*(\boldsymbol{\pi}) \\ &= \arg \min F(\mathbf{r}/\tau) - \mathbf{r} \cdot \boldsymbol{\pi}/\tau + F^*(\boldsymbol{\pi}) \\ &= \arg \min \mathbf{KL}(\boldsymbol{\pi} \parallel \mathbf{p}) \text{ where } \mathbf{p} = e^{\mathbf{r}/\tau - F(\mathbf{r}/\tau)} \end{aligned}$$

**Suggests a natural surrogate**

$$\begin{aligned} & \arg \min \mathbf{KL}(\mathbf{p} \parallel \boldsymbol{\pi}) = \arg \min F(\mathbf{q}) - \mathbf{q} \cdot \mathbf{p} \\ & \bullet \text{ convex in } \mathbf{q} \end{aligned}$$

Optimize policy  $\pi : X \rightarrow \Delta^n$

$$\pi(a \mid x) = e^{q(x)_a - F(q(x))}$$

$$F(q(x)) = \log \sum_a e^{q(x)_a}$$

$$q : X \rightarrow \mathbb{R}^n \quad \text{neural network}$$

$$\text{Let } F^*(\boldsymbol{\pi}) = \boldsymbol{\pi} \cdot \log \boldsymbol{\pi}$$

# Optimization objectives

Back to **general** rewards

	$a_1$	$a_2$	$\dots$	$a_n$
$x_1$	$r_{11}$	$r_{12}$	$r_{1\dots}$	$r_{1n}$
$x_2$	$r_{21}$	$r_{22}$	$r_{2\dots}$	$r_{2n}$
$x_3$	$r_{31}$	$r_{32}$	$r_{3\dots}$	$r_{3n}$
$x_4$	$r_{41}$	$r_{42}$	$r_{4\dots}$	$r_{4n}$
$x_5$	$r_{51}$	$r_{52}$	$r_{5\dots}$	$r_{5n}$
$x_6$	$r_{61}$	$r_{62}$	$r_{6\dots}$	$r_{6n}$
$\vdots$	$r_{:1}$	$r_{:2}$	$r_{: \dots}$	$r_{:n}$
$x_m$	$r_{m1}$	$r_{m2}$	$r_{m\dots}$	$r_{mn}$

## Comparison to maximum likelihood

before  $-\mathbf{r} \cdot \log \pi = F(\mathbf{q}) - \mathbf{q} \cdot \mathbf{r}$

now  $\mathbf{KL}(\mathbf{p} \parallel \pi) \equiv F(\mathbf{q}) - \mathbf{q} \cdot \mathbf{p}$

## If $\mathbf{r} = \mathbf{1}_a$ is an indicator

- become equivalent as  $\tau \rightarrow 0$
- $\lim_{\tau \rightarrow 0} \mathbf{p} = \mathbf{r} = \mathbf{1}_a$
- but  $\tau > 0$  gives **soft targets** for **KL**  
"label smoothing"

improves generalization in practice

Optimize policy  $\pi : X \rightarrow \Delta^n$

$$\pi(a | x) = e^{q(x)_a - F(q(x))}$$

$$F(q(x)) = \log \sum_a e^{q(x)_a}$$

$q : X \rightarrow \mathbb{R}^n$  neural network

# Optimization objectives

Back to **general** rewards

	$a_1$	$a_2$	$\dots$	$a_n$
$x_1$	$r_{11}$	$r_{12}$	$r_{1\dots}$	$r_{1n}$
$x_2$	$r_{21}$	$r_{22}$	$r_{2\dots}$	$r_{2n}$
$x_3$	$r_{31}$	$r_{32}$	$r_{3\dots}$	$r_{3n}$
$x_4$	$r_{41}$	$r_{42}$	$r_{4\dots}$	$r_{4n}$
$x_5$	$r_{51}$	$r_{52}$	$r_{5\dots}$	$r_{5n}$
$x_6$	$r_{61}$	$r_{62}$	$r_{6\dots}$	$r_{6n}$
$\vdots$	$r_{:1}$	$r_{:2}$	$r_{: \dots}$	$r_{:n}$
$x_m$	$r_{m1}$	$r_{m2}$	$r_{m\dots}$	$r_{mn}$

**A convex, calibrated upper bound**

$$\mathbf{KL}(\pi \| \mathbf{p}) \leq \mathbf{KL}(\mathbf{p} \| \pi) + \frac{\tau}{4} \|\mathbf{r}/\tau - \mathbf{q}\|^2$$

Optimize policy  $\pi : X \rightarrow \Delta^n$


$$\pi(a | x) = e^{q(x)_a - F(q(x))}$$

$$F(q(x)) = \log \sum_a e^{q(x)_a}$$

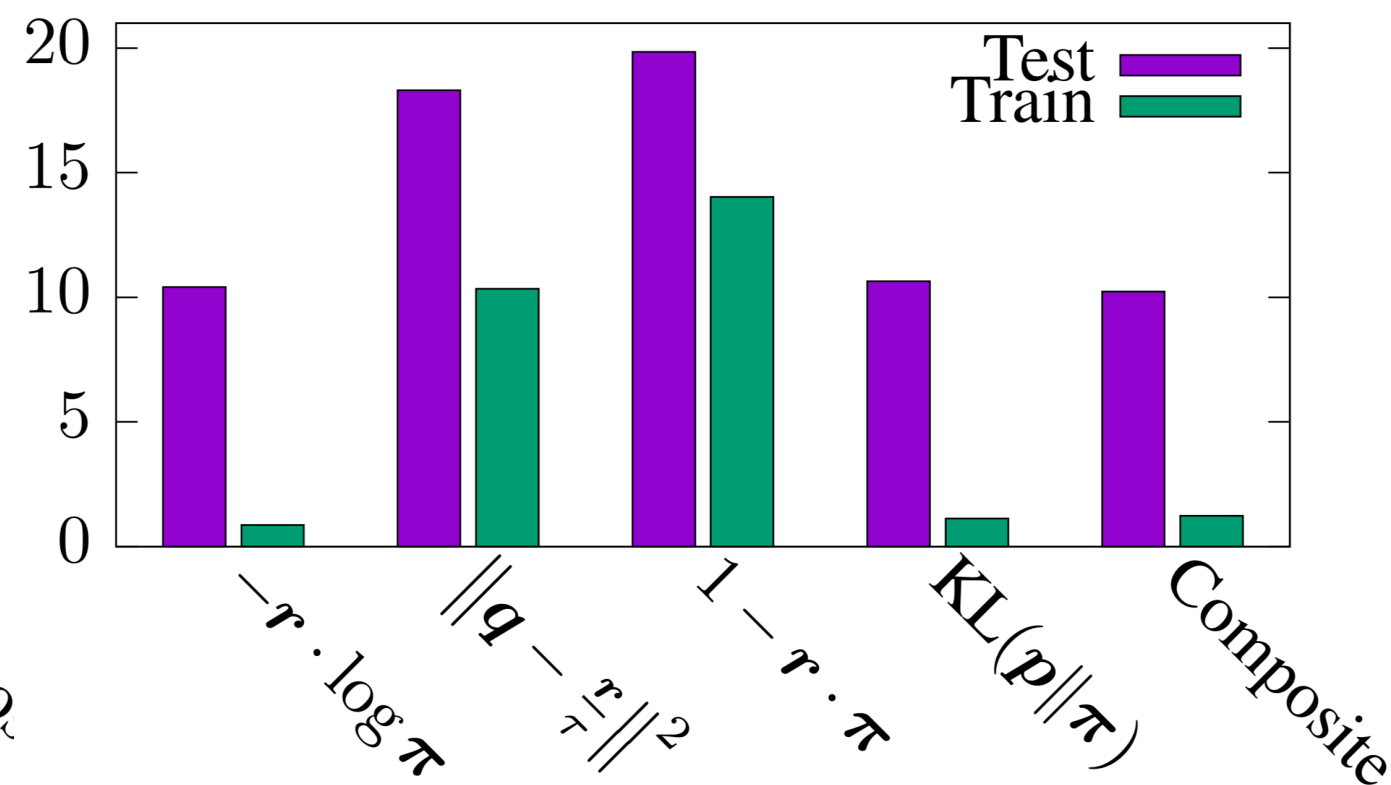
$$q : X \rightarrow \mathfrak{R}^n \quad \text{neural network}$$

# Optimization objectives

MNIST



CIFAR10



# Batch policy optimization

	$a_1$	$a_2$	$\dots$	$a_n$
$x_1$	$r_{11}$	$r_{12}$	$r_{1\dots}$	$r_{1n}$
$x_2$	$r_{21}$	$r_{22}$	$r_{2\dots}$	$r_{2n}$
$x_3$	$r_{31}$	$r_{32}$	$r_{3\dots}$	$r_{3n}$
$x_4$	$r_{41}$	$r_{42}$	$r_{4\dots}$	$r_{4n}$
$x_5$	$r_{51}$	$r_{52}$	$r_{5\dots}$	$r_{5n}$
$x_6$	$r_{61}$	$r_{62}$	$r_{6\dots}$	$r_{6n}$
$\vdots$	$r_{:1}$	$r_{:2}$	$r_{:....}$	$r_{:n}$
$x_m$	$r_{m1}$	$r_{m2}$	$r_{m\dots}$	$r_{mn}$

## Three key issues

1. generalization

training objective  
 $\neq$

2. optimization

target objective

3. missing data

# Batch policy optimization

	$a_1$	$a_2$	...	$a_n$
$x_1$		$r_1$		
$x_2$				$r_2$
$x_3$	$r_3$			
$x_4$			$r_4$	
$x_5$	$r_5$			
$x_6$		$r_6$		
$\vdots$			$r_i$	
$x_m$				$r_m$


## Three key issues

1. generalization
2. optimization
3. missing data

# Supervised vs reinforcement learning

supervised classification

	$a_1$	$a_2$	...	$a_n$
$x_1$	$r_{11}$	$r_{12}$	$r_{1...}$	$r_{1n}$
$x_2$	$r_{21}$	$r_{22}$	$r_{2...}$	$r_{2n}$
$x_3$	$r_{31}$	$r_{32}$	$r_{3...}$	$r_{3n}$
$x_4$	$r_{41}$	$r_{42}$	$r_{4...}$	$r_{4n}$
$x_5$	$r_{51}$	$r_{52}$	$r_{5...}$	$r_{5n}$
$x_6$	$r_{61}$	$r_{62}$	$r_{6...}$	$r_{6n}$
$\vdots$	$r_{:1}$	$r_{:2}$	$r_{:....}$	$r_{:n}$
$x_m$	$r_{m1}$	$r_{m2}$	$r_{m...}$	$r_{mn}$



batch policy optimization

	$a_1$	$a_2$	...	$a_n$
$x_1$		$r_1$		
$x_2$				$r_2$
$x_3$	$r_3$			
$x_4$			$r_4$	
$x_5$	$r_5$			
$x_6$		$r_6$		
$\vdots$			$r_{:}$	
$x_m$				$r_m$

Optimize policy  $\pi : X \rightarrow \Delta^n$  to maximize expected reward on **test** contexts

**key difference is missing data**



# Missing data inference

How to handle missing data?

	$a_1$	$a_2$	...	$a_n$
$x_1$		$r_1$		
$x_2$				$r_2$
$x_3$	$r_3$			
$x_4$			$r_4$	
$x_5$	$r_5$			
$x_6$		$r_6$		
$\vdots$			$r_i$	
$x_m$				$r_m$

Optimize policy  $\pi : X \rightarrow \Delta^n$

# Missing data inference

	$a_1$	$a_2$	...	$a_n$
$x_1$	$q_{11}$	$r_1$	$q_{1...}$	$q_{1n}$
$x_2$	$q_{21}$	$q_{22}$	$q_{2...}$	$r_2$
$x_3$	$r_3$	$q_{32}$	$q_{3...}$	$q_{3n}$
$x_4$	$q_{41}$	$q_{42}$	$r_4$	$q_{4n}$
$x_5$	$r_5$	$q_{52}$	$q_{5...}$	$q_{5n}$
$x_6$	$q_{61}$	$r_6$	$q_{6...}$	$q_{6n}$
$\vdots$	$q_{:1}$	$q_{:2}$	$r_{:}$	$q_{:n}$
$x_m$	$q_{m1}$	$q_{m2}$	$q_{m...}$	$r_m$

## Simple idea

### imputation

- fill in guesses for missing values
- reduce to fully observed case

## Might sound naive

- but this is actually a dominant approach

Optimize policy  $\pi : X \rightarrow \Delta^n$

# Missing data inference

	$a_1$	$a_2$	...	$a_n$
$x_1$	0	$\frac{r_1}{\beta_1}$	0	0
$x_2$	0	0	0	$\frac{r_2}{\beta_2}$
$x_3$	$\frac{r_3}{\beta_3}$	0	0	0
$x_4$	0	0	$\frac{r_4}{\beta_4}$	0
$x_5$	$\frac{r_5}{\beta_5}$	0	0	0
$x_6$	0	$\frac{r_6}{\beta_6}$	0	0
$\vdots$	0	0	$\frac{r_i}{\beta_i}$	0
$x_m$	0	0	0	$\frac{r_m}{\beta_m}$

Optimize policy  $\pi : X \rightarrow \Delta^n$

## Example

**importance corrected expected reward**

$$\max \sum_i \frac{\pi(a_i | x_i)}{\beta_i} r_i$$

where  $\beta$  are proposal probabilities from behavior strategy

We already know this is a poor objective but what about missing data inference?

**Equivalent** to  $\max \hat{\mathbf{r}} \cdot \boldsymbol{\pi}$  using

$$\hat{\mathbf{r}}_i = \mathbf{1}_{a_i} \frac{r_i}{\beta_i}$$

**That is**

- **exaggerate** observed values by  $1/\beta_i$
- fill in all unobserved values with **0**

# Missing data inference

	$a_1$	$a_2$	...	$a_n$
$x_1$	0	$\frac{r_1}{\beta_1}$	0	0
$x_2$	0	0	0	$\frac{r_2}{\beta_2}$
$x_3$	$\frac{r_3}{\beta_3}$	0	0	0
$x_4$	0	0	$\frac{r_4}{\beta_4}$	0
$x_5$	$\frac{r_5}{\beta_5}$	0	0	0
$x_6$	0	$\frac{r_6}{\beta_6}$	0	0
$\vdots$	0	0	$\frac{r_{\cdot}}{\beta_{\cdot}}$	0
$x_m$	0	0	0	$\frac{r_m}{\beta_m}$

**This is a pretty lame inference principle**

- **altering** the data we do see
- to compensate for a bad guess about the data we don't see

**But ... its unbiased!**

$$\mathbb{E}[\hat{\mathbf{r}} | x] = \sum_a \beta_a \mathbf{1}_a \frac{r_a}{\beta_a} = \sum_a \mathbf{1}_a r_a = \mathbf{r}$$

Optimize policy  $\pi : X \rightarrow \Delta^n$

# Missing data inference

	$a_1$	$a_2$	...	$a_n$
$x_1$	$\tau q_{11}$	$\lambda(r_1 - \tau q_{12})$	$\tau q_{1...}$	$\tau q_{1n}$
$x_2$	$\tau q_{21}$	$\tau q_{22}$	$\tau q_{2...}$	$\lambda(r_2 - \tau q_{2n})$
$x_3$	$\lambda(r_3 - \tau q_{31})$	$\tau q_{32}$	$\tau q_{3...}$	$\tau q_{3n}$
$x_4$	$\tau q_{41}$	$\tau q_{42}$	$\lambda(r_4 - \tau q_{4...})$	$\tau q_{4n}$
$x_5$	$\lambda(r_5 - \tau q_{51})$	$\tau q_{52}$	$\tau q_{5...}$	$\tau q_{5n}$
$x_6$	$\tau q_{61}$	$\lambda(r_6 - \tau q_{62})$	$\tau q_{6...}$	$\tau q_{6n}$
$\vdots$	$\tau q_{:1}$	$\tau q_{:2}$	$\lambda(r_{:} - \tau q_{:...})$	$\tau q_{:n}$
$x_m$	$\tau q_{m1}$	$\tau q_{m2}$	$\tau q_{m...}$	$\lambda(r_m - \tau q_{mn})$

Optimize policy  $\pi : X \rightarrow \Delta^n$

## Improvement

“doubly robust estimation”

- instead of filling in with 0s
- fill in with guesses from a model  $\mathbf{q}(x)$

$$\hat{\mathbf{r}} = \tau \mathbf{q} + \lambda \mathbf{1}_a (r - \tau q_a)$$

## Also unbiased

- as long as  $\lambda = 1/\beta_i$
- but still alters observed data

## Where should the model come from?

- could use a separate critic
- train via least squares, then optimize  $\pi$
- works okay, but not great

## Note

- there is only one action value function for single-step decision making,  $r(x, a)$
- actor-critic approaches trivialized

# Missing data inference

	$a_1$	$a_2$	...	$a_n$
$x_1$	$\tau q_{11}$	$\lambda(r_1 - \tau q_{12})$	$\tau q_{1...}$	$\tau q_{1n}$
$x_2$	$\tau q_{21}$	$\tau q_{22}$	$\tau q_{2...}$	$\lambda(r_2 - \tau q_{2n})$
$x_3$	$\lambda(r_3 - \tau q_{31})$	$\tau q_{32}$	$\tau q_{3...}$	$\tau q_{3n}$
$x_4$	$\tau q_{41}$	$\tau q_{42}$	$\lambda(r_4 - \tau q_{4...})$	$\tau q_{4n}$
$x_5$	$\lambda(r_5 - \tau q_{51})$	$\tau q_{52}$	$\tau q_{5...}$	$\tau q_{5n}$
$x_6$	$\tau q_{61}$	$\lambda(r_6 - \tau q_{62})$	$\tau q_{6...}$	$\tau q_{6n}$
$\vdots$	$\tau q_{:1}$	$\tau q_{:2}$	$\lambda(r_{:} - \tau q_{:...})$	$\tau q_{:n}$
$x_m$	$\tau q_{m1}$	$\tau q_{m2}$	$\tau q_{m...}$	$\lambda(r_m - \tau q_{mn})$

Optimize policy  $\pi : X \rightarrow \Delta^n$

## Unified approach

- actor and critic are same model
- $\pi = e^{\mathbf{q} - F(\mathbf{q})}$  where  $F(\mathbf{q}) = \log \mathbf{1} \cdot e^{\mathbf{q}}$
- use logits  $\tau \mathbf{q}(x)$  to predict rewards

$$q(x, a) \approx \frac{r(x, a)}{\tau}$$


## Can combine with previous objectives

- $\text{KL}(\pi \| \hat{\mathbf{p}})$  where  $\hat{\mathbf{p}} = e^{\hat{\mathbf{r}}/\tau - F(\hat{\mathbf{r}}/\tau)}$
- $\text{KL}(\hat{\mathbf{p}} \| \pi)$
- $\text{KL}(\pi \| \hat{\mathbf{p}}) \leq \text{KL}(\hat{\mathbf{p}} \| \pi) + \frac{\tau}{4} \|\hat{\mathbf{r}}/\tau - \mathbf{q}\|^2$


these are somewhat sensitive to ranking, unlike least squares

# Missing data inference


MNIST



CIFAR10



# Missing data inference



## Even more principled approach

- back to first principles
- how do we reason about missing data in the rest of ML and statistics?

Optimize policy  $\pi : X \rightarrow \text{exp-family}(\mathfrak{R})$


$$\pi(a | x) = e^{q(x)_a - F(q(x))}$$

$$F(q(x)) = \log \int e^{q(x)_a} \mu(da)$$

$q : X \rightarrow \mathfrak{R}^k$  neural network



# Missing data inference



Optimize policy  $\pi : X \rightarrow \text{exp-family}(\mathfrak{R})$   
 $\pi(a | x) = e^{q(x)_a - F(q(x))}$   
 $F(q(x)) = \log \int e^{q(x)_a} \mu(da)$   
 $q : X \rightarrow \mathfrak{R}^k$  neural network

## Even more principled approach

- back to first principles
- how do we reason about missing data in the rest of ML and statistics?

## Bayesian inference

- postulate a generative model of reward

$$q \rightarrow \xi \rightarrow r$$

- **e.g. Gaussian**

- prior  $\xi \sim \mathcal{N}(q, Q)$
- likelihood  $r | a, \xi \sim \mathcal{N}(\phi(a) \cdot \xi, \sigma^2)$
- posterior  $\xi | r_0, a_0 \sim \mathcal{N}(\mu, C)$

$$\mu = C(\phi(a_0)r_0\sigma^{-2} + Q^{-1}q)$$


$$C = (Q^{-1} + \sigma^{-2}\phi(a_0)\phi(a_0)^\top)^{-1}$$

- predictive  $r | a, r_0, a_0 \sim \mathcal{N}(\hat{\mu}, \hat{\sigma}^2)$

$$\hat{\mu} = \phi(a) \cdot \mu$$

$$\hat{\sigma}^2 = \sigma^2 + \phi(a_0)^\top C \phi(a_0)$$

# Missing data inference



Optimize policy  $\pi : X \rightarrow \text{exp-family}(\mathfrak{R})$

$$\pi(a | x) = e^{q(x)_a - F(q(x))}$$

$$F(q(x)) = \log \int e^{q(x)_a} \mu(da)$$

$q : X \rightarrow \mathfrak{R}^k$  neural network

## Empirical Bayes estimation

- optimize hyperparameters  $\mathbf{q}$  (neural network)
- integrate out parameters  $\xi$

## Example

marginal likelihood

$$\begin{aligned} & -\log p(r_0 | a_0, \mathbf{q}) \\ &= -\log \int p(r_0 | a_0, \xi) p(\xi | \mathbf{q}) d\xi \\ &= \frac{1}{2\sigma^2} (\phi(a_0) \cdot q - r_0)^2 + \frac{1}{2} \log \sigma^2 + c \end{aligned}$$

- essentially least squares regression


## Can alternatively use surrogates

$\min \mathbf{KL}(\text{prior} || \text{posterior})$

$\min \mathbf{KL}(\text{posterior} || \text{prior}) \approx \min I(\xi; r_0)$

# Missing data inference

Sum of squared test error on continuous action MNIST ( $a \in \mathcal{R}^{10}$ )



# Batch policy optimization

	$a_1$	$a_2$	...	$a_n$
$x_1$		$r_1$		
$x_2$				$r_2$
$x_3$	$r_3$			
$x_4$			$r_4$	
$x_5$	$r_5$			
$x_6$		$r_6$		
$\vdots$			$r_i$	
$x_m$				$r_m$

## Three key issues

1. generalization

training objective  
 $\neq$


2. optimization

target objective

3. missing data

classical methods  
still help

# The RL problem



1. multi-agent interaction → non-stationarity
2. partial observability → construct memory
3. exploration → explore/exploit
4. sequential decisions → temporal credit assignment
5. exploitation → policy optimization

## 4. Fundamentals of sequential decision making

approximate dynamic programming

- be afraid, be very afraid

policy optimization with simple value estimates

- surrogate objectives and missing data inference show promise

RL in the dual

- avoid the deadly triad
- new approaches to MCMC stationary distribution inference